

Building a Student Intervention System

Supervised Learning Project

Template code

Open the template iPython notebook `student_intervention.ipynb` and follow along.

Project brief

As education has grown to rely more and more on technology, more and more data is available for examination and prediction. Logs of student activities, grades, interactions with teachers and fellow students, and more are now captured through learning management systems like Canvas and Edmodo and available in real time. This is especially true for online classrooms, which are becoming more and more popular even at the middle and high school levels.

Within all levels of education, there exists a push to help increase the likelihood of student success without watering down the education or engaging in behaviors that raise the likelihood of passing metrics without improving the actual underlying learning. Graduation rates are often the criteria of choice for this, and educators and administrators are after new ways to predict success and failure early enough to stage

effective interventions, as well as to identify the effectiveness of different interventions.

Toward that end, your goal as a software engineer hired by the local school district is to model the factors that predict how likely a student is to pass their high school final exam. The school district has a goal to reach a 95% graduation rate by the end of the decade by identifying students who need intervention before they drop out of school. You being a clever engineer decide to implement a student intervention system using concepts you learned from supervised machine learning. Instead of buying expensive servers or implementing new data models from the ground up, you reach out to a 3rd party company who can provide you the necessary software libraries and servers to run your software.

However, with limited resources and budgets, the board of supervisors wants you to find the most effective model with the least amount of computation costs (you pay the company by the memory and CPU time you use on their servers). In order to build the intervention software, you first will need to analyze the dataset on students' performance. Your goal is to choose and develop a model that will predict the likelihood that a given student will pass, thus helping diagnose whether or not an intervention is necessary. Your model must be developed based on a subset of the data that we provide to you, and it will be tested against a subset of the data that is kept hidden from the learning algorithm, in order to test the model's effectiveness on data outside the training set.

Deliverables

1. Classification vs Regression

This is a classification problem. Because the predicted value is 0 or 1. We need to classify students into two categories instead of predicting a numerical value. So obviously, this is a classification problem.

2. Exploring the Data

By preprocessing the data, we could see these facts:

- Total number of students: 395
- Number of students who passed: 265
- Number of students who failed: 130
- Number of features: 30
- Graduation rate of the class: 67.00%

3. Preparing the Data

```
Feature column(s):-
```

```
['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu',  
'Fedu', 'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime',  
'studytime', 'failures', 'schoolsup', 'famsup', 'paid',  
'activities', 'nursery', 'higher', 'internet', 'romantic',
```

```
'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences']
```

Target column: passed

The appendix section contains information on each of these attributes.

4. Training and Evaluating Models

Choose 3 supervised learning models that are available in scikit-learn, and appropriate for this problem?

I choose the following four models to analyze the dataset:

- Decision Tree
- Random Forest
- Support Vector Machine

Note: In the ipython notebook decision trees was used. However, it was used merely for a benchmark on f1 score.

Decision Tree

Description:

- Decision tree is a classification algorithms which would be great for classify 0 or 1 problem.

Complexity Analysis:

- It has a space complexity of $O(n)$ where n is the number of instances.
- The training time complexity of decision tree model is $O(m \cdot n \log n)$, where n is the number of the number of instances and m is the number of features.

With the figures above, we should see that the expected training and prediction time for NBC should be less due to its linear nature.

Pros:

- It is simple to understand and to interpret. Trees can be visualised.

Cons:

- It is easy to get overfitted. And optimal algorithm could be NP-hard, which makes it very unstable.

Reasons for Selection:

- Decision tree is a good model to predict 0/1 output by splitting at different features, which looks very suitable for this problem.

Training set size	100	200	296
Training time (secs)	0.001	0.005	0.001
Prediction time (secs)	0.000	0.002	0.000

F1 score for training set	1.000	1.000	1.000
F1 score for test set	0.550	0.688	0.677

Thus, for the entire training data, the decision tree will give a F1 score of 0.677 for testing set. So obviously it has a serious overfitting.

Random Forest

Description:

- A Random Forest has a space complexity of $O(m^{(1/2)} * n \log n)$, where m is the number of features and n the number of elements in the dataset, under the assumption that a reasonably symmetric tree is built.
- The training complexity is given as $O(M * m^{(1/2)} * n \log n)$, where M denotes the number of trees.

Pros:

- Random forest could deal with unbalanced data and handle missing values well.

Cons:

- Random Forest may also have overfitting problem when there are too many missing values.

Reasons for Selection:

- This algorithm was chosen as it is an ensemble of decision tree classifiers, which will dismiss the affection caused by unuseful features.

Training set size	100	200	296
Training time (secs)	0.006	0.006	0.018
Prediction time (secs)	0.001	0.001	0.003
F1 score for training set	0.993	0.993	0.992
F1 score for test set	0.712	0.750	0.765

The classifier was not tuned and default parameters used. Random Forest also have a overfitting issue.

Support Vector Machine

Description:

- In machine learning, support vector machines are supervised learning models with associated learning algorithms that analyze data and recognize patterns. It is very good for 0/1 output learning problems.

Complexity Analysis:

- SVM has a space complexity of $O(n^2)$
- It has a training time of $O(n^3)$ where n is the training dataset size.

Pros:

- Effective in high dimensional spaces.
- Still effective in cases where number of dimensions is greater than the number of samples.

Cons:

- It has a high computational cost. Also SVM is also sensitive to noise.

Reasons for Selection:

- I expected the data to not be linear with very high dimension. So by using an appropriate kernel like rbf, we would expect to effectively tune the classifier for a great f1 score.

Training set size	100	200	296
Training time (secs)	0.002	0.005	0.010
Prediction time (secs)	0.001	0.002	0.003
F1 score for training set	0.839	0.876	0.880
F1 score for test set	0.765	0.758	0.757

5. Choosing the Best Model

From result of section 4, I would choose Support vector machine provides the best performance. It has a great balance between training

time and F1 score, also it does not have a overfitting issue.

Support Vector Machines will give a line based on training data and try to divide the training data into two parts. The fitted line will find the maximum distance to the nearest point of both sizes. In this specific case, the fitted line will predict if a student will graduate or not. Then the line will try to draw a boundary between them.

The chosen SVM model was tuned using Grid Search due to the size of data. Also, in such a case where the data is unbalanced, so I would choose F1 for metric in GridSearch. The parameters optimized were `gamma`, `C` and `tolerance`. And I choose rbf as the kernel.

The below is a output from IPython note book. And F1 score is 0.768 for testing data.

```
Successfully fit a model!
```

```
SVC(C=200, cache_size=200, class_weight=None, coef0=0.0, degree=3,  
    gamma=0.001, kernel='rbf', max_iter=-1, probability=False,  
    random_state=None, shrinking=True, tol=0.001, verbose=False)
```

```
Predicting labels using SVC...
```

```
Done!
```

```
Predicting labels using SVC...
```

```
Done!
```

```
Prediction time (secs): 0.005
```

F1 score for training set: 0.884875846501

Predicting labels using SVC...

Done!

Prediction time (secs): 0.001

F1 score for test set: 0.768115942029

Data

The dataset used in this project is included as `student-data.csv`. This dataset has the following attributes:

- `school` ? student's school (binary: "GP" or "MS")
- `sex` ? student's sex (binary: "F" - female or "M" - male)
- `age` ? student's age (numeric: from 15 to 22)
- `address` ? student's home address type (binary: "U" - urban or "R" - rural)
- `famsize` ? family size (binary: "LE3" - less or equal to 3 or "GT3" - greater than 3)
- `Pstatus` ? parent's cohabitation status (binary: "T" - living together or "A" - apart)
- `Medu` ? mother's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - "5th to 9th grade", 3 - secondary education or 4 - "higher education")
- `Fedu` ? father's education (numeric: 0 - none, 1 - primary education (4th grade), 2 - 5th to 9th grade, 3 - secondary education or 4 - "higher education")
- `Mjob` ? mother's job (nominal: "teacher", "health" care

related, civil "services" (e.g. administrative or police),
"at_home" or "other")

- `Fjob` ? father's job (nominal: "teacher", "health" care
related, civil "services" (e.g. administrative or police),
"at_home" or "other")

- `reason` ? reason to choose this school (nominal: close t
o "home", school "reputation", "course" preference or "othe
r")

- `guardian` ? student's guardian (nominal: "mother", "fath
er" or "other")

- `traveltime` ? home to school travel time (numeric: 1 - <
15 min., 2 - 15 to 30 min., 3 - 30 min. to 1 hour, or 4 - >
1 hour)

- `studytime` ? weekly study time (numeric: 1 - <2 hours, 2
- 2 to 5 hours, 3 - 5 to 10 hours, or 4 - >10 hours)

- `failures` ? number of past class failures (numeric: n if
1<=n<3, else 4)

- `schoolsup` ? extra educational support (binary: yes or n
o)

- `famsup` ? family educational support (binary: yes or no)

- `paid` ? extra paid classes within the course subject (Ma
th or Portuguese) (binary: yes or no)

- `activities` ? extra-curricular activities (binary: yes o
r no)

- `nursery` ? attended nursery school (binary: yes or no)

- `higher` ? wants to take higher education (binary: yes or
no)

- `internet` ? Internet access at home (binary: yes or no)

- `romantic` ? with a romantic relationship (binary: yes or no)
- `famrel` ? quality of family relationships (numeric: from 1 - very bad to 5 - excellent)
- `freetime` ? free time after school (numeric: from 1 - very low to 5 - very high)
- `goout` ? going out with friends (numeric: from 1 - very low to 5 - very high)
- `Dalc` ? workday alcohol consumption (numeric: from 1 - very low to 5 - very high)
- `Walc` ? weekend alcohol consumption (numeric: from 1 - very low to 5 - very high)
- `health` ? current health status (numeric: from 1 - very bad to 5 - very good)
- `absences` ? number of school absences (numeric: from 0 to 93)
- `passed` ? did the student pass the final exam (binary: yes or no)